

# AI Driven Fraud Detection System for Financial Transactions

This project focuses on building an AI-driven fraud detection system for financial transactions using machine learning. The main goal is to accurately identify fraudulent transactions within a highly imbalanced, anonymised dataset.

The key steps followed in this notebook are:

- Exploratory Data Analysis (EDA) to understand class imbalance
- Preprocessing of important features ('Amount' and 'Time') to improve model learning
- Handling severe class imbalance using oversampling techniques (SMOTE and ADASYN)
- Training and evaluating multiple machine learning models (Random Forest, Logistic Regression, SVM, KNN)
- Comparing models based on Recall, F1 Score, and Precision-Recall AUC to select the best-performing pipeline

Throughout the notebook, a strong focus is placed on:

- Handling imbalanced data effectively
- Maintaining reproducibility and fairness
- Using clear metrics that reflect real-world fraud detection needs

The results from this project provide insights into how different models and resampling techniques affect fraud detection performance in a practical setting and which is best.

```
# Import necessary libraries for data analysis and visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Machine Learning libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import RobustScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
```

```

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix, roc_auc_score, roc_curve
from sklearn.metrics import classification_report
from sklearn.metrics import roc_auc_score
from sklearn.metrics import average_precision_score
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import auc

# Resampling techniques for imbalanced dataset
from imblearn.over_sampling import SMOTE
from imblearn.over_sampling import ADASYN
from collections import Counter

# For visualizations
sns.set(style="whitegrid")

import json # For Saving and loading metrics
import joblib # For saving and loading models
import pickle # For saving and loading metrics
import os # For checking file existences
from scipy.stats import boxcox # For box cox transformations

# Mount Google Drive to access the file
from google.colab import drive
drive.mount('/content/drive')

# Define the file path to your dataset in Google Drive
file_path = '/content/drive/My Drive/Colab Notebooks/creditcard.csv'

# Load the dataset
df = pd.read_csv(file_path)

# Display the first few rows of the dataset
df.head()

Mounted at /content/drive

{"type": "dataframe", "variable_name": "df"}

df.describe()

{"type": "dataframe"}

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -

```

```
0    Time    284807 non-null float64
1     V1     284807 non-null float64
2     V2     284807 non-null float64
3     V3     284807 non-null float64
4     V4     284807 non-null float64
5     V5     284807 non-null float64
6     V6     284807 non-null float64
7     V7     284807 non-null float64
8     V8     284807 non-null float64
9     V9     284807 non-null float64
10    V10    284807 non-null float64
11    V11    284807 non-null float64
12    V12    284807 non-null float64
13    V13    284807 non-null float64
14    V14    284807 non-null float64
15    V15    284807 non-null float64
16    V16    284807 non-null float64
17    V17    284807 non-null float64
18    V18    284807 non-null float64
19    V19    284807 non-null float64
20    V20    284807 non-null float64
21    V21    284807 non-null float64
22    V22    284807 non-null float64
23    V23    284807 non-null float64
24    V24    284807 non-null float64
25    V25    284807 non-null float64
26    V26    284807 non-null float64
27    V27    284807 non-null float64
28    V28    284807 non-null float64
29    Amount 284807 non-null float64
30    Class  284807 non-null int64
```

```
dtypes: float64(30), int64(1)
```

```
memory usage: 67.4 MB
```

```
# Check for null values
```

```
print("Missing values per column:\n", df.isnull().sum())
```

```
Missing values per column:
```

```
Time    0
V1       0
V2       0
V3       0
V4       0
V5       0
V6       0
V7       0
V8       0
V9       0
V10      0
V11      0
```

```
V12      0
V13      0
V14      0
V15      0
V16      0
V17      0
V18      0
V19      0
V20      0
V21      0
V22      0
V23      0
V24      0
V25      0
V26      0
V27      0
V28      0
Amount    0
Class     0
dtype: int64
```

```
#Finding all rows where fraud (Class = 1)
frauds = df.loc[df['Class'] == 1]
#Finding all rows that are not fraud (Class = 0)
non_frauds = df.loc[df['Class'] == 0]

#Calculating the counts and percentages
fraud_count = len(frauds)
non_fraud_count = len(non_frauds)

#Calculating percentages
fraud_percentage = round(fraud_count / len(df) * 100, 2)
non_fraud_percentage = round(non_fraud_count / len(df) * 100, 2)

#Printing the results
print('Fraud - ', fraud_count, 'transactions or ', fraud_percentage,
      '% of the dataset')
print('Legitimate - ', non_fraud_count, 'transactions or ',
      non_fraud_percentage, '% of the dataset')

Fraud - 492 transactions or 0.17 % of the dataset
Legitimate - 284315 transactions or 99.83 % of the dataset
```

### **#Train/Test Split and Baseline**

```
from sklearn.model_selection import train_test_split

# Define features (X) and target (y)
X = df.drop(columns=['Class']) # drop target column
y = df['Class'] # define target
```

```

# Train-test split with stratification for class balance
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

# Display split details
print(f"Training set shape: {X_train.shape}, Target: {y_train.shape}")
print(f"Testing set shape: {X_test.shape}, Target: {y_test.shape}")
print(f"\nFraud cases in training set: {y_train.sum()}")
print(f"Fraud cases in testing set: {y_test.sum()}")

Training set shape: (227845, 30), Target: (227845,)
Testing set shape: (56962, 30), Target: (56962,)

Fraud cases in training set: 394
Fraud cases in testing set: 98

```

We split the original dataset into a training set (80%) and a testing set (20%) using stratified sampling to maintain the class distribution. This ensures that the minority fraud cases are proportionally represented in both sets, avoiding skewed evaluation results later.

Random Forest

```

# File path for Random Forest metrics
rf_path = "/content/drive/My Drive/Colab
Notebooks/rf_metrics_RAW.json"

# Load or train Random Forest
if os.path.exists(rf_path):
    print("Loaded saved Random Forest classification report:")
    with open(rf_path, "r") as f:
        saved_metrics = json.load(f)
    print(pd.DataFrame(saved_metrics).transpose())
else:
    print("No saved classification report found. Training Random
Forest model on RAW data...")
    rf_model = RandomForestClassifier(random_state=42)
    rf_model.fit(X_train, y_train)
    y_pred_rf = rf_model.predict(X_test)
    report = classification_report(y_test, y_pred_rf,
output_dict=True)
    with open(rf_path, "w") as f:
        json.dump(report, f, indent=4)
    print(pd.DataFrame(report).transpose())

Loaded saved Random Forest classification report:

```

|   | precision | recall   | f1-score | support      |
|---|-----------|----------|----------|--------------|
| 0 | 0.999684  | 0.999912 | 0.999798 | 56864.000000 |
| 1 | 0.941176  | 0.816327 | 0.874317 | 98.000000    |

|              |          |          |          |              |
|--------------|----------|----------|----------|--------------|
| accuracy     | 0.999596 | 0.999596 | 0.999596 | 0.999596     |
| macro avg    | 0.970430 | 0.908119 | 0.937057 | 56962.000000 |
| weighted avg | 0.999583 | 0.999596 | 0.999582 | 56962.000000 |

## Logistic Regression

```
# File path for Logistic Regression metrics
lr_path = "/content/drive/My Drive/Colab
Notebooks/lr_metrics_RAW.json"

# Load or train Logistic Regression
if os.path.exists(lr_path):
    print("Loaded saved Logistic Regression classification report:")
    with open(lr_path, "r") as f:
        saved_metrics = json.load(f)
    print(pd.DataFrame(saved_metrics).transpose())
else:
    print("No saved classification report found. Training Logistic
Regression model on RAW data...")
    lr_model = LogisticRegression(max_iter=1000, random_state=42)
    lr_model.fit(X_train, y_train)
    y_pred_lr = lr_model.predict(X_test)
    report = classification_report(y_test, y_pred_lr,
output_dict=True)
    with open(lr_path, "w") as f:
        json.dump(report, f, indent=4)
    print(pd.DataFrame(report).transpose())
```

```
Loaded saved Logistic Regression classification report:
              precision    recall  f1-score   support
0               0.999437    0.999771    0.999604    56864.00000
1               0.835443    0.673469    0.745763     98.00000
accuracy               0.999210    0.999210    0.999210     0.99921
macro avg               0.917440    0.836620    0.872684    56962.00000
weighted avg               0.999155    0.999210    0.999168    56962.00000
```

## KNN

```
# File path for KNN metrics
knn_path = "/content/drive/My Drive/Colab
Notebooks/knn_metrics_RAW.json"

# Load or train KNN
if os.path.exists(knn_path):
    print("Loaded saved KNN classification report:")
    with open(knn_path, "r") as f:
        saved_metrics = json.load(f)
    print(pd.DataFrame(saved_metrics).transpose())
else:
```

```

    print("No saved classification report found. Training KNN model on
RAW data...")
    knn_model = KNeighborsClassifier()
    knn_model.fit(X_train, y_train)
    y_pred_knn = knn_model.predict(X_test)
    report = classification_report(y_test, y_pred_knn,
output_dict=True)
    with open(knn_path, "w") as f:
        json.dump(report, f, indent=4)
    print(pd.DataFrame(report).transpose())

```

Loaded saved KNN classification report:

|              | precision | recall   | f1-score | support      |
|--------------|-----------|----------|----------|--------------|
| 0            | 0.998332  | 1.000000 | 0.999165 | 56864.000000 |
| 1            | 1.000000  | 0.030612 | 0.059406 | 98.000000    |
| accuracy     | 0.998332  | 0.998332 | 0.998332 | 0.998332     |
| macro avg    | 0.999166  | 0.515306 | 0.529286 | 56962.000000 |
| weighted avg | 0.998335  | 0.998332 | 0.997549 | 56962.000000 |

SVM

```

# File path for SVM metrics
svm_path = "/content/drive/My Drive/Colab
Notebooks/svm_metrics_RAW.json"

# Load or train SVM
if os.path.exists(svm_path):
    print("Loaded saved SVM classification report:")
    with open(svm_path, "r") as f:
        saved_metrics = json.load(f)
    print(pd.DataFrame(saved_metrics).transpose())
else:
    print("No saved classification report found. Training SVM model on
RAW data...")
    svm_model = SVC(kernel='rbf', random_state=42)
    svm_model.fit(X_train, y_train)
    y_pred_svm = svm_model.predict(X_test)
    report = classification_report(y_test, y_pred_svm,
output_dict=True)
    with open(svm_path, "w") as f:
        json.dump(report, f, indent=4)
    print(pd.DataFrame(report).transpose())

```

Loaded saved SVM classification report:

|          | precision | recall   | f1-score | support      |
|----------|-----------|----------|----------|--------------|
| 0        | 0.998280  | 1.000000 | 0.999139 | 56864.000000 |
| 1        | 0.000000  | 0.000000 | 0.000000 | 98.000000    |
| accuracy | 0.998280  | 0.998280 | 0.998280 | 0.998280     |

|              |          |          |          |              |
|--------------|----------|----------|----------|--------------|
| macro avg    | 0.499140 | 0.500000 | 0.499570 | 56962.000000 |
| weighted avg | 0.996562 | 0.99828  | 0.997420 | 56962.000000 |

## #Exploratory Data Analysis

Lets have a look at how many of the transactions are fraudulent and how many are legitimate.

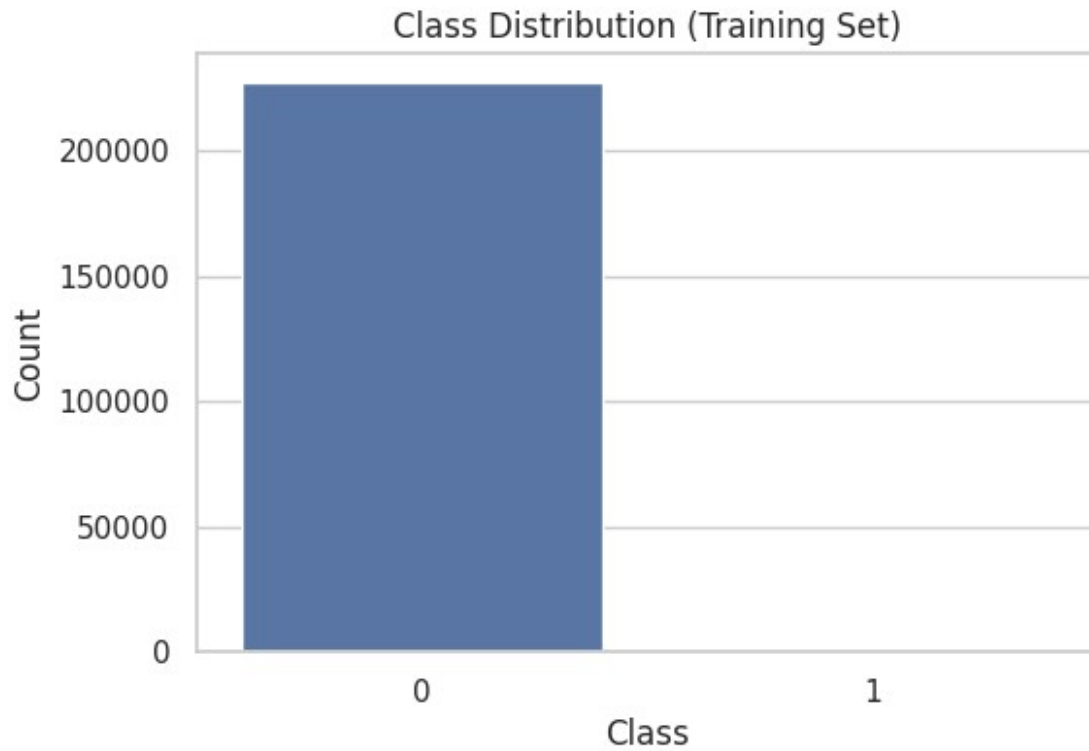
```
# 1. Check class distribution in the training set (clear format)
print("Class Distribution in Training Set:\n")
class_counts = y_train.value_counts()
class_percentages = y_train.value_counts(normalize=True) * 100
distribution_df = pd.DataFrame({
    "Count": class_counts,
    "Percentage": class_percentages.round(4)
})
distribution_df.index = distribution_df.index.map({0: 'Legitimate
(0)', 1: 'Fraudulent (1)'})
display(distribution_df)
```

```
# Bar plot of class distribution
plt.figure(figsize=(6,4))
sns.countplot(x=y_train)
plt.title('Class Distribution (Training Set)')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()
```

Class Distribution in Training Set:

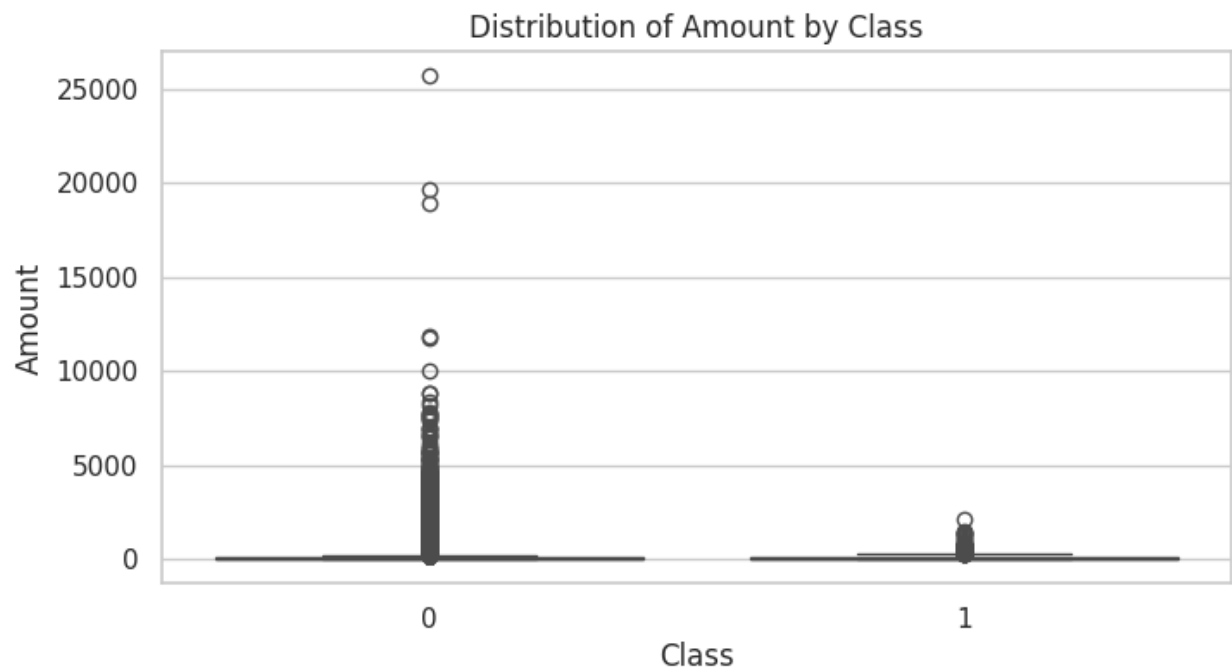
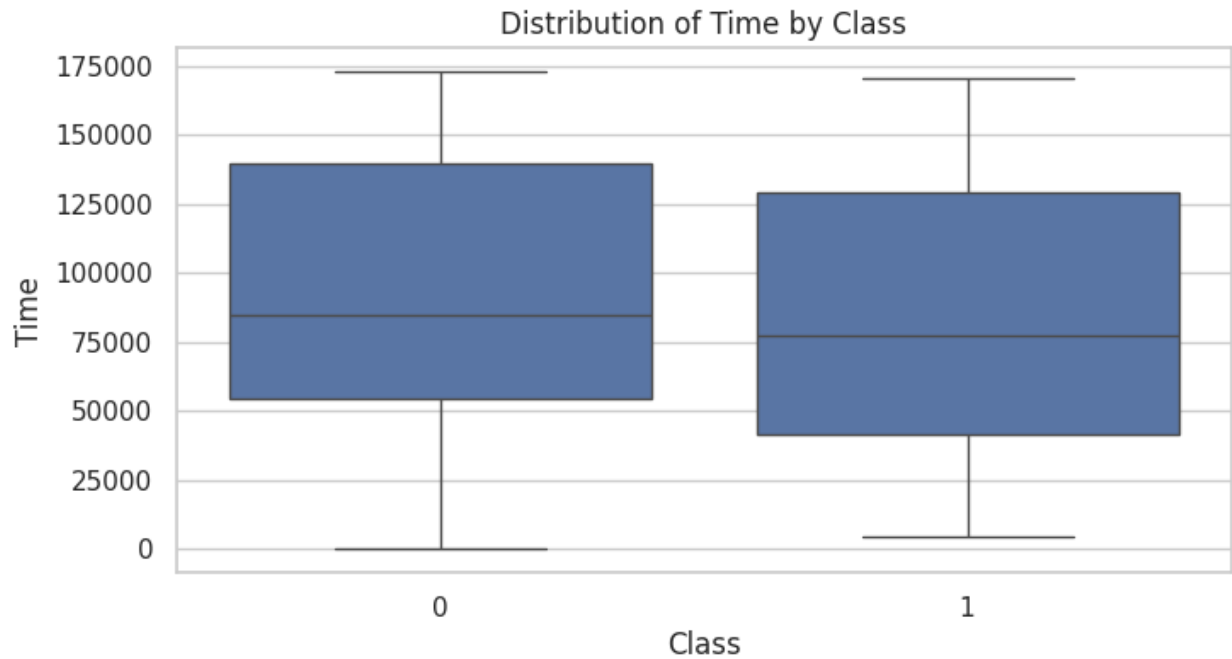
```
{"summary":{"\n  \"name\": \"distribution_df\",\n  \"rows\": 2,\n  \"fields\": [\n    {\n      \"column\": \"Class\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"Fraudulent (1)\",\n          \"Legitimate (0)\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"Count\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 160553,\n          \"min\": 394,\n          \"max\": 227451,\n          \"num_unique_values\": 2,\n          \"samples\": [\n            394,\n            227451\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\": \"Percentage\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 70.46616059372045,\n            \"min\": 0.1729,\n            \"max\": 99.8271,\n            \"num_unique_values\": 2,\n            \"samples\": [\n              0.1729,\n              99.8271\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          }\n        }\n      ]\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"distribution_df\"}
```





We checked the class distribution in the training set. As expected, fraud cases are extremely rare compared to non-fraud cases, highlighting the severe imbalance problem that needs to be addressed before model training.


```
for col in ['Time', 'Amount']:
    plt.figure(figsize=(8, 4))
    sns.boxplot(x=y_train, y=X_train[col])
    plt.title(f'Distribution of {col} by Class')
    plt.show()
```



```
# Correlation matrix on training set (for all features + class)
df_train_full = X_train.copy()
df_train_full['Class'] = y_train
corr = df_train_full.corr()

# Plot correlation heatmap
plt.figure(figsize=(20, 18))
sns.heatmap(corr, annot=False, cmap='coolwarm', center=0)
```

```
plt.title('Correlation Matrix (Training Set)')
plt.show()
```



We plotted the correlation matrix to explore relationships between the available features. Although the dataset is PCA-transformed (making feature meanings unknown), the correlation matrix still helps identify any slight patterns or dependencies. Understanding feature correlations can guide model expectations — for example, highly correlated features could lead to redundant information, while uncorrelated features suggest independent contributions.

```
#Top 10 positively/negatively correlated features with Class
class_corr = corr['Class'].drop('Class')
top_positive = class_corr.sort_values(ascending=False).head(10)
```

```

top_negative = class_corr.sort_values().head(10)

print("\nTop 10 Positively Correlated Features with 'Class':")
print(top_positive)
print("\nTop 10 Negatively Correlated Features with 'Class':")
print(top_negative)

```

Top 10 Positively Correlated Features with 'Class':

```

V11      0.153709
V4       0.135014
V2       0.090586
V21      0.035588
V19      0.032380
V8       0.020552
V20      0.019385
V27      0.016034
V28      0.009810
Amount   0.006211

```

Name: Class, dtype: float64

Top 10 Negatively Correlated Features with 'Class':

```

V17     -0.321937
V14     -0.301054
V12     -0.259989
V10     -0.217894
V3      -0.194135
V16     -0.193826
V7      -0.186184
V18     -0.108732
V1      -0.100041
V9      -0.098247

```


Name: Class, dtype: float64

## Data Preprocessing

```

plt.figure(figsize=(10, 2))
sns.boxplot(x=X_train['Amount'], orient='h', color='lightblue')
plt.title("Box Plot of Transaction Amount (Training Set)")
plt.xlabel("Amount")
plt.grid(True)
plt.tight_layout()
plt.show()

```



```
#Shift amount to avoid zero (Box-Cox needs strictly positive values)
X_train['Amount_shifted'] = X_train['Amount'] + 1e-9
X_test['Amount_shifted'] = X_test['Amount'] + 1e-9
```

```
#Fit Box-Cox on training set only
X_train['scaled_amount'], fitted_lambda =
boxcox(X_train['Amount_shifted'])
```

```
# Save the fitted lambda for Box-Cox transformation
joblib.dump(fitted_lambda, 'boxcox_lambda_amount.pkl')
```

```
# Load existing saved Box-Cox lambda (.pkl file)
boxcox_lambda = joblib.load('boxcox_lambda_amount.pkl')
```

```
# Save the lambda into a JSON file
with open('boxcox_lambda_amount.json', 'w') as f:
    json.dump({'lambda': float(boxcox_lambda)}, f)
```

```
print("Saved Box-Cox lambda as boxcox_lambda_amount.json")
```

```
# Apply the same transformation using the learned lambda to the test
set
```

```
X_test['scaled_amount'] = boxcox(X_test['Amount_shifted'],
lambda=fitted_lambda)
```

```
# Check skewness before transformation
original_skew = X_train['Amount'].skew()
print(f"Skewness of original 'Amount': {original_skew:.4f}")
```


```
# Check skewness after Box-Cox transformation
transformed_skew = pd.Series(X_train['scaled_amount']).skew()
print(f"Skewness of Box-Cox transformed 'Amount':
{transformed_skew:.4f}")
```

```
# Drop the original and intermediate columns
X_train.drop(['Amount', 'Amount_shifted'], axis=1, inplace=True)
X_test.drop(['Amount', 'Amount_shifted'], axis=1, inplace=True)
```

```
plt.figure(figsize=(12, 2))
sns.boxplot(x=X_train['scaled_amount'], orient='h')
plt.title("Box Plot of Scaled Transaction Amount (Box-Cox
Transformed)")
```

```
plt.xlabel("Scaled Amount")
plt.grid(True)
plt.show()
```

Saved Box-Cox lambda as boxcox\_lambda\_amount.json  
 Skewness of original 'Amount': 18.1939  
 Skewness of Box-Cox transformed 'Amount': 0.1146




We applied Box-Cox transformation to the 'Amount' feature to reduce its skewness. Normalizing the feature helps machine learning models perform better because highly skewed data can bias the model towards certain feature values.

Now lets move onto time

```
# Convert time to hours for better readability
X_train_time_dist = X_train.copy()
X_train_time_dist['Hour'] = (X_train_time_dist['Time'] // 3600) % 24

plt.figure(figsize=(10, 4))
sns.histplot(X_train_time_dist['Hour'], bins=24, kde=False)
plt.title("Distribution of Transactions Over Hour of Day (Original Time Feature)")
plt.xlabel("Hour of Day")
plt.ylabel("Number of Transactions")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
# Copy original datasets
X_train_time = X_train.copy()
X_test_time = X_test.copy()

# Extract Hour from 'Time' column (seconds since first transaction)
X_train_time['Hour'] = (X_train_time['Time'] // 3600) % 24
X_test_time['Hour'] = (X_test_time['Time'] // 3600) % 24

# Apply sin/cos transformation to encode cyclical nature
X_train_time['Hour_sin'] = np.sin(2 * np.pi * X_train_time['Hour'] / 24)
X_train_time['Hour_cos'] = np.cos(2 * np.pi * X_train_time['Hour'] / 24)

X_test_time['Hour_sin'] = np.sin(2 * np.pi * X_test_time['Hour'] / 24)
X_test_time['Hour_cos'] = np.cos(2 * np.pi * X_test_time['Hour'] / 24)

# Drop raw 'Time' column and 'Hour'
X_train_time.drop(columns=['Time', 'Hour'], inplace=True)
X_test_time.drop(columns=['Time', 'Hour'], inplace=True)
```

We engineered new 'Hour\_sin' and 'Hour\_cos' features from the 'Time' column using sine and cosine transformations. This approach captures the cyclic nature of time (e.g., transactions occurring at similar times of day) which could be useful for detecting fraud patterns.

```
# Merge the updated X_train_time with y_train for correlation analysis
train_df = X_train_time.copy()
train_df['Class'] = y_train.values

# Compute correlation matrix
corr_matrix = train_df.corr()

# Plot correlation heatmap
plt.figure(figsize=(16, 14))
```


```
sns.heatmap(corr_matrix, cmap='coolwarm', annot=False, center=0,
square=True, linewidths=0.3)
plt.title(" Correlation Matrix After Feature Engineering (Training
Set)")
plt.show()

# Show top correlated features with 'Class'
corr_with_target = corr_matrix['Class'].drop('Class')

top_positive_corr =
corr_with_target.sort_values(ascending=False).head(10)
top_negative_corr = corr_with_target.sort_values().head(10)

print("Top 10 Positively Correlated Features with 'Class':")
print(top_positive_corr)
print("\nTop 10 Negatively Correlated Features with 'Class':")
print(top_negative_corr)
```





#### Top 10 Positively Correlated Features with 'Class':

|          |          |
|----------|----------|
| V11      | 0.153709 |
| V4       | 0.135014 |
| V2       | 0.090586 |
| V21      | 0.035588 |
| V19      | 0.032380 |
| V8       | 0.020552 |
| V20      | 0.019385 |
| V27      | 0.016034 |
| Hour_sin | 0.011816 |
| V28      | 0.009810 |

Name: Class, dtype: float64

#### Top 10 Negatively Correlated Features with 'Class':

|     |           |
|-----|-----------|
| V17 | -0.321937 |
|-----|-----------|

```
V14    -0.301054
V12    -0.259989
V10    -0.217894
V3      -0.194135
V16    -0.193826
V7      -0.186184
V18    -0.108732
V1      -0.100041
V9      -0.098247
Name: Class, dtype: float64
```

## Pre Processed Baseline

```
# Check structure and basic stats of preprocessed training set
print("\n Structure of Preprocessed Training Set:")
X_train_time.info()

print("\n Descriptive Statistics of Preprocessed Features:")
display(X_train_time.describe())

print("\n Sample of Preprocessed Training Data:")
display(X_train_time.head())
```

```
Structure of Preprocessed Training Set:
<class 'pandas.core.frame.DataFrame'>
Index: 227845 entries, 265518 to 17677
Data columns (total 31 columns):
#   Column          Non-Null Count  Dtype
---  -
0   V1               227845 non-null float64
1   V2               227845 non-null float64
2   V3               227845 non-null float64
3   V4               227845 non-null float64
4   V5               227845 non-null float64
5   V6               227845 non-null float64
6   V7               227845 non-null float64
7   V8               227845 non-null float64
8   V9               227845 non-null float64
9   V10              227845 non-null float64
10  V11              227845 non-null float64
11  V12              227845 non-null float64
12  V13              227845 non-null float64
13  V14              227845 non-null float64
14  V15              227845 non-null float64
15  V16              227845 non-null float64
16  V17              227845 non-null float64
17  V18              227845 non-null float64
```

|    |               |        |          |         |
|----|---------------|--------|----------|---------|
| 18 | V19           | 227845 | non-null | float64 |
| 19 | V20           | 227845 | non-null | float64 |
| 20 | V21           | 227845 | non-null | float64 |
| 21 | V22           | 227845 | non-null | float64 |
| 22 | V23           | 227845 | non-null | float64 |
| 23 | V24           | 227845 | non-null | float64 |
| 24 | V25           | 227845 | non-null | float64 |
| 25 | V26           | 227845 | non-null | float64 |
| 26 | V27           | 227845 | non-null | float64 |
| 27 | V28           | 227845 | non-null | float64 |
| 28 | scaled_amount | 227845 | non-null | float64 |
| 29 | Hour_sin      | 227845 | non-null | float64 |
| 30 | Hour_cos      | 227845 | non-null | float64 |

dtypes: float64(31)  
memory usage: 55.6 MB

Descriptive Statistics of Preprocessed Features:

```
{"type": "dataframe"}
```

Sample of Preprocessed Training Data:

```
{"type": "dataframe"}
```

Before we Preprocess the data and create the splits, we will run it on the raw dataset to set the baseline

Random Forest Model

```
import os, json
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# File path for Random Forest metrics
rf_path = "/content/drive/My Drive/Colab Notebooks/rf_metrics.json"

# Function to save classification report
def save_classification_report(model_name, y_true, y_pred, path):
    report = classification_report(y_true, y_pred, output_dict=True)
    with open(path, "w") as f:
        json.dump(report, f, indent=4)
    print(f"{model_name} classification report saved to {path}")

# Load or train Random Forest
if os.path.exists(rf_path):
    print("Loaded saved Random Forest classification report:")
    with open(rf_path, "r") as f:
        saved_metrics = json.load(f)
    df_metrics_rf = pd.DataFrame(saved_metrics).transpose()
```

```

    print(df_metrics_rf)
else:
    print("No saved classification report found. Training Random
Forest model...")
    rf_model = RandomForestClassifier(random_state=42)
    rf_model.fit(X_train_time, y_train)
    y_pred_rf = rf_model.predict(X_test_time)
    save_classification_report("Random Forest", y_test, y_pred_rf,
rf_path)

```

Loaded saved Random Forest classification report:

|              | precision | recall   | f1-score | support      |
|--------------|-----------|----------|----------|--------------|
| 0            | 0.999701  | 0.999930 | 0.999815 | 56864.000000 |
| 1            | 0.952941  | 0.826531 | 0.885246 | 98.000000    |
| accuracy     | 0.999631  | 0.999631 | 0.999631 | 0.999631     |
| macro avg    | 0.976321  | 0.913230 | 0.942531 | 56962.000000 |
| weighted avg | 0.999621  | 0.999631 | 0.999618 | 56962.000000 |

Logisitic Regression

```

from sklearn.linear_model import LogisticRegression

# File path for Logistic Regression metrics
lr_path = "/content/drive/My Drive/Colab Notebooks/lr_metrics.json"

# Load or train Logistic Regression
if os.path.exists(lr_path):
    print("Loaded saved Logistic Regression classification report:")
    with open(lr_path, "r") as f:
        saved_metrics = json.load(f)
    df_metrics_lr = pd.DataFrame(saved_metrics).transpose()
    print(df_metrics_lr)
else:
    print("No saved classification report found. Training Logistic
Regression model...")
    lr_model = LogisticRegression(max_iter=1000, random_state=42)
    lr_model.fit(X_train_time, y_train)
    y_pred_lr = lr_model.predict(X_test_time)
    save_classification_report("Logistic Regression", y_test,
y_pred_lr, lr_path)

```

Loaded saved Logistic Regression classification report:

|              | precision | recall   | f1-score | support      |
|--------------|-----------|----------|----------|--------------|
| 0            | 0.999402  | 0.999754 | 0.999578 | 56864.000000 |
| 1            | 0.820513  | 0.653061 | 0.727273 | 98.000000    |
| accuracy     | 0.999157  | 0.999157 | 0.999157 | 0.999157     |
| macro avg    | 0.909958  | 0.826408 | 0.863425 | 56962.000000 |
| weighted avg | 0.999095  | 0.999157 | 0.999110 | 56962.000000 |

KNN Model

```

from sklearn.neighbors import KNeighborsClassifier

# File path for KNN metrics
knn_path = "/content/drive/My Drive/Colab Notebooks/knn_metrics.json"

# Load or train KNN
if os.path.exists(knn_path):
    print("Loaded saved KNN classification report:")
    with open(knn_path, "r") as f:
        saved_metrics = json.load(f)
        df_metrics_knn = pd.DataFrame(saved_metrics).transpose()
        print(df_metrics_knn)
else:
    print("No saved classification report found. Training KNN model...")
    knn_model = KNeighborsClassifier()
    knn_model.fit(X_train_time, y_train)
    y_pred_knn = knn_model.predict(X_test_time)
    save_classification_report("KNN", y_test, y_pred_knn, knn_path)

```

Loaded saved KNN classification report:

|              | precision | recall   | f1-score | support      |
|--------------|-----------|----------|----------|--------------|
| 0            | 0.999631  | 0.999842 | 0.999736 | 56864.000000 |
| 1            | 0.895349  | 0.785714 | 0.836957 | 98.000000    |
| accuracy     | 0.999473  | 0.999473 | 0.999473 | 0.999473     |
| macro avg    | 0.947490  | 0.892778 | 0.918346 | 56962.000000 |
| weighted avg | 0.999451  | 0.999473 | 0.999456 | 56962.000000 |

SVM

```

from sklearn.svm import SVC

# File path for SVM metrics
svm_path = "/content/drive/My Drive/Colab Notebooks/svm_metrics.json"

# Load or train SVM
if os.path.exists(svm_path):
    print("Loaded saved SVM classification report:")
    with open(svm_path, "r") as f:
        saved_metrics = json.load(f)
        df_metrics_svm = pd.DataFrame(saved_metrics).transpose()
        print(df_metrics_svm)
else:
    print("No saved classification report found. Training SVM model...")
    svm_model = SVC(random_state=42)
    svm_model.fit(X_train_time, y_train)
    y_pred_svm = svm_model.predict(X_test_time)
    save_classification_report("SVM", y_test, y_pred_svm, svm_path)

```

Loaded saved SVM classification report:

|              | precision | recall   | f1-score | support      |
|--------------|-----------|----------|----------|--------------|
| 0            | 0.999543  | 0.999947 | 0.999745 | 56864.000000 |
| 1            | 0.960000  | 0.734694 | 0.832370 | 98.000000    |
| accuracy     | 0.999491  | 0.999491 | 0.999491 | 0.999491     |
| macro avg    | 0.979771  | 0.867321 | 0.916057 | 56962.000000 |
| weighted avg | 0.999475  | 0.999491 | 0.999457 | 56962.000000 |

Comparing the baseline results

```
import matplotlib.pyplot as plt
import numpy as np

# Preprocessed metrics for Class 1 (fraud)
metrics = {
    "Random Forest": {"precision": 0.952941, "recall": 0.826531, "f1-score": 0.885246},
    "Logistic Regression": {"precision": 0.820513, "recall": 0.653061, "f1-score": 0.727273},
    "KNN": {"precision": 0.895349, "recall": 0.785714, "f1-score": 0.836957},
    "SVM": {"precision": 0.960000, "recall": 0.734694, "f1-score": 0.832370},
}

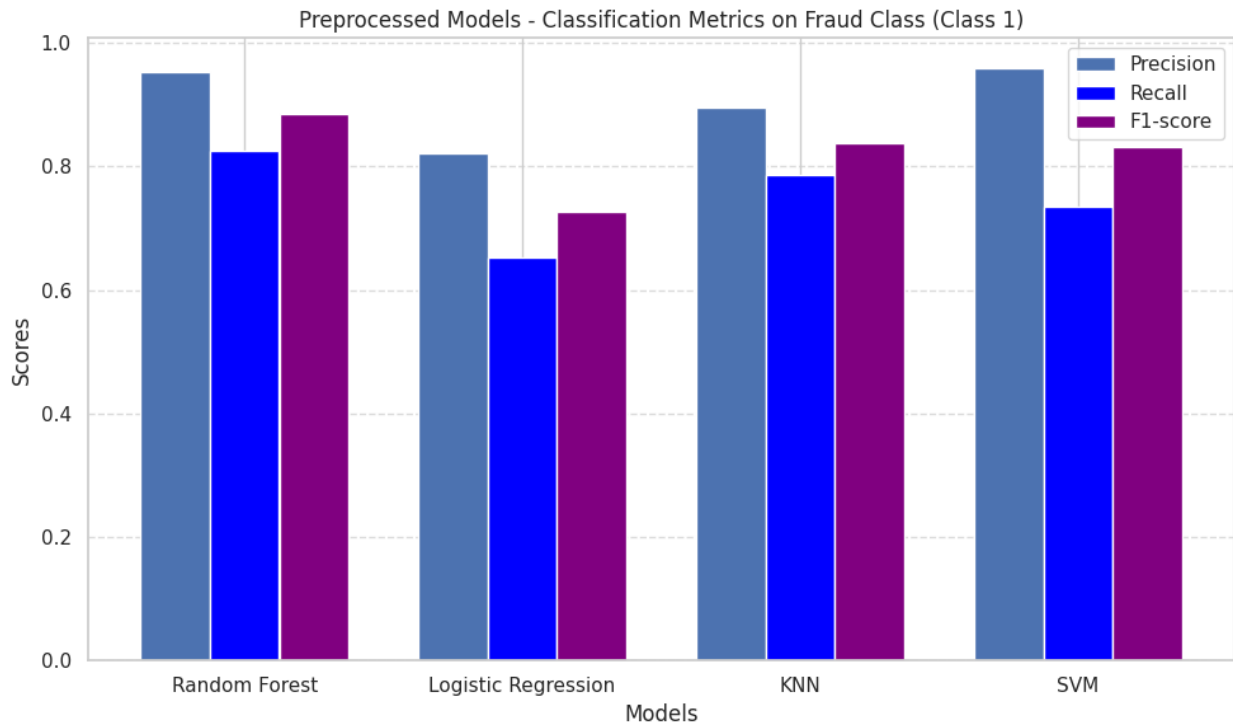
models = list(metrics.keys())
precision = [metrics[m]["precision"] for m in models]
recall = [metrics[m]["recall"] for m in models]
f1_score = [metrics[m]["f1-score"] for m in models]

x = np.arange(len(models))
width = 0.25

fig, ax = plt.subplots(figsize=(10, 6))
bars1 = ax.bar(x - width, precision, width, label="Precision")
bars2 = ax.bar(x, recall, width, label="Recall", color="blue")
bars3 = ax.bar(x + width, f1_score, width, label="F1-score", color="purple")

ax.set_xlabel("Models")
ax.set_ylabel("Scores")
ax.set_title("Preprocessed Models - Classification Metrics on Fraud Class (Class 1)")
ax.set_xticks(x)
ax.set_xticklabels(models)
ax.legend()
ax.grid(axis="y", linestyle="--", alpha=0.7)

plt.tight_layout()
plt.show()
```



## Data Preprocessing (SMOTE/ADASYN)

```
# Before SMOTE
print("Before SMOTE:", Counter(y_train))

# Initialize SMOTE
smote = SMOTE(random_state=42)

# Apply SMOTE to preprocessed training data
X_train_smote, y_train_smote = smote.fit_resample(X_train_time,
y_train)

# After SMOTE
print("After SMOTE:", Counter(y_train_smote))

Before SMOTE: Counter({0: 227451, 1: 394})
After SMOTE: Counter({0: 227451, 1: 227451})
```

We used SMOTE (Synthetic Minority Over-sampling Technique) to balance the training data by generating synthetic fraud examples. This reduces the bias toward the majority class and improves the model's ability to correctly identify fraudulent transactions.

```
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt


# Run t-SNE with fixed random state
```

```

tsne_smote = TSNE(n_components=2, random_state=42, perplexity=30)
X_smote_tsne = tsne_smote.fit_transform(X_train_smote)

# Plot
plt.figure(figsize=(8, 6))
plt.scatter(X_smote_tsne[y_train_smote == 0, 0],
            X_smote_tsne[y_train_smote == 0, 1],
            label='Non-Fraud', alpha=0.5, s=10, c='skyblue')
plt.scatter(X_smote_tsne[y_train_smote == 1, 0],
            X_smote_tsne[y_train_smote == 1, 1],
            label='Fraud', alpha=0.7, s=10, c='red')
plt.title("t-SNE Visualization of SMOTE-Resampled Training Data")
plt.legend()
plt.grid(True)
plt.show()

```



```

# Before ADASYN
print("Before ADASYN:", Counter(y_train))

# Initialize ADASYN

```



```

adasyn = ADASYN(random_state=42)

# Apply ADASYN to preprocessed training data
X_train_adasyn, y_train_adasyn = adasyn.fit_resample(X_train_time,
y_train)

# After ADASYN
print("After ADASYN:", Counter(y_train_adasyn))

Before ADASYN: Counter({0: 227451, 1: 394})
After ADASYN: Counter({1: 227460, 0: 227451})

```


We also applied ADASYN (Adaptive Synthetic Sampling) to generate more synthetic minority examples, focusing more on harder-to-classify cases. This provides another approach to balance the dataset and allows us to compare results between SMOTE and ADASYN later.

```

# Run t-SNE with fixed random state
tsne_adasyn = TSNE(n_components=2, random_state=42, perplexity=30)
X_adasyn_tsne = tsne_adasyn.fit_transform(X_train_adasyn)

# Plot
plt.figure(figsize=(8, 6))
plt.scatter(X_adasyn_tsne[y_train_adasyn == 0, 0],
X_adasyn_tsne[y_train_adasyn == 0, 1],
            label='Non-Fraud', alpha=0.5, s=10, c='skyblue')
plt.scatter(X_adasyn_tsne[y_train_adasyn == 1, 0],
X_adasyn_tsne[y_train_adasyn == 1, 1],
            label='Fraud', alpha=0.7, s=10, c='red')
plt.title("t-SNE Visualization of ADASYN-Resampled Training Data")
plt.legend()
plt.grid(True)
plt.show()

```



## Model Training

##SMOTE

```
import json
import os
import pandas as pd
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    roc_auc_score,
    precision_recall_curve,
    auc
)
import matplotlib.pyplot as plt
import seaborn as sns

# Dictionary to store AUC-ROC and PR AUC scores
internal_auc_scores = {}
```

```

# Helper function to save and display results
def save_metrics_and_outputs(model_name, y_true, y_pred, y_proba,
report_path, internal_auc_dict):
    # === Save and display classification report ===
    report = classification_report(y_true, y_pred, output_dict=True)
    with open(report_path, 'w') as f:
        json.dump(report, f, indent=4)
    print(f"\n{model_name} Classification Report:")
    display(pd.DataFrame(report).transpose())

    # === Display confusion matrix ===
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(6, 5))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=[0,
1], yticklabels=[0, 1])
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(f"{model_name} Confusion Matrix")
    plt.show()

    # === Calculate and store AUC-ROC and PR AUC ===
    auc_roc = roc_auc_score(y_true, y_proba[:, 1])
    precision, recall, _ = precision_recall_curve(y_true, y_proba[:,
1])
    pr_auc = auc(recall, precision)
    internal_auc_dict[model_name] = {
        "AUC_ROC": auc_roc,
        "PR_AUC": pr_auc
    }
    print(f"{model_name} AUC-ROC: {auc_roc:.4f}, PR AUC: {pr_auc:.4f}
(Saved internally)")

```

Random Forest

```

rf_smote_path = "/content/drive/My Drive/Colab
Notebooks/rf_SMOTE_metrics.json"

if os.path.exists(rf_smote_path):
    with open(rf_smote_path, 'r') as f:
        saved_report = json.load(f)
    print("Random Forest Classification Report:")
    display(pd.DataFrame(saved_report).transpose())

    rf_model = RandomForestClassifier(random_state=42)
    rf_model.fit(X_train_smote, y_train_smote)
    y_pred_rf = rf_model.predict(X_test_time)
    y_proba_rf = rf_model.predict_proba(X_test_time)
    save_metrics_and_outputs("Random Forest", y_test, y_pred_rf,

```

```
y_proba_rf, rf_smote_path, internal_auc_scores)
```

```
else:
```


```
    print("No saved report found. Training Random Forest on SMOTE  
data...")
```

```
    rf_model = RandomForestClassifier(random_state=42)  
    rf_model.fit(X_train_smote, y_train_smote)  
    y_pred_rf = rf_model.predict(X_test_time)  
    y_proba_rf = rf_model.predict_proba(X_test_time)  
    save_metrics_and_outputs("Random Forest", y_test, y_pred_rf,  
y_proba_rf, rf_smote_path, internal_auc_scores)
```

No saved report found. Training Random Forest on SMOTE data...

Random Forest Classification Report:

```
{"summary": "{\n  \"name\": \"    save_metrics_and_outputs(\\\"\\\"\\\"Random  
Forest\\\"\\\"\\\", y_test, y_pred_rf, y_proba_rf, rf_smote_path,  
internal_auc_scores)\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"precision\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.04502686328713626, \n        \"min\": 0.898876404494382, \n        \"max\": 0.9996835053540344, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.898876404494382, \n          0.9995100722603184, \n          0.9995259997893332\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\", \n        \"column\": \"precision\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 0.08196470676738973, \n          \"min\": 0.8163265306122449, \n          \"max\": 0.9998417276308385, \n          \"num_unique_values\": 4, \n          \"samples\": [\n            0.8163265306122449, \n            0.9080841291215417, \n            0.9998417276308385\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\", \n          \"column\": \"f1-score\", \n          \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.06438359730356975, \n            \"min\": 0.8556149732620321, \n            \"max\": 0.9997626102323782, \n            \"num_unique_values\": 5, \n            \"samples\": [\n              0.8556149732620321, \n              0.9995146121209514, \n              0.9995259997893332\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\", \n            \"column\": \"support\", \n            \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 31154.412478809652, \n              \"min\": 0.9995259997893332, \n              \"max\": 56962.0, \n              \"num_unique_values\": 4, \n              \"samples\": [\n                98.0, \n                56962.0, \n                56864.0\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\" \n            } \n          } \n        ] \n      } \n    } \n  ], \n  \"type\": \"dataframe\"}
```



Random Forest AUC-ROC: 0.9683, PR AUC: 0.8684 (Saved internally)

Logistic regression

```
lr_smote_path = "/content/drive/My Drive/Colab  
Notebooks/lr_SMOTE_metrics.json"  
  
if os.path.exists(lr_smote_path):  
    with open(lr_smote_path, 'r') as f:  
        saved_report = json.load(f)  
    print("Logistic Regression Classification Report:")  
    display(pd.DataFrame(saved_report).transpose())  
  
    lr_model = LogisticRegression(max_iter=1000, random_state=42)  
    lr_model.fit(X_train_smote, y_train_smote)  
    y_pred_lr = lr_model.predict(X_test_time)  
    y_proba_lr = lr_model.predict_proba(X_test_time)  
    save_metrics_and_outputs("Logistic Regression", y_test, y_pred_lr,  
y_proba_lr, lr_smote_path, internal_auc_scores)  
else:
```

```

print("No saved report found. Training Logistic Regression on
SMOTE data...")
lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train_smote, y_train_smote)
y_pred_lr = lr_model.predict(X_test_time)
y_proba_lr = lr_model.predict_proba(X_test_time)
save_metrics_and_outputs("Logistic Regression", y_test, y_pred_lr,
y_proba_lr, lr_smote_path, internal_auc_scores)

```


No saved report found. Training Logistic Regression on SMOTE data...

Logistic Regression Classification Report:

```

{"summary":{"\n  \"name\": \"      save_metrics_and_outputs(\\\"\\\"\\\"Logistic
Regression\\\"\\\", y_test, y_pred_lr, y_proba_lr, lr_smote_path,
internal_auc_scores)\\\", \"rows\": 5, \"fields\": [\n    {\n      \"column\": \"precision\", \"properties\": {\n        \"dtype\": \"number\", \"std\": 0.41817983689130067, \"n
        \"min\": 0.05415162454873646, \"max\": 0.9998553345388789, \"n
        \"num_unique_values\": 5, \"samples\": [\n
          0.05415162454873646, \"          0.9982283031218108, \"n
          0.9722622098943156\n          ], \"semantic_type\": \"\", \"n
          \"description\": \"\"\n          }, \"column\":
          \"recall\", \"properties\": {\n            \"dtype\": \"number\", \"n
            \"std\": 0.02411293361417645, \"min\": 0.9183673469387755, \"n
            \"max\": 0.9723550928531233, \"num_unique_values\": 4, \"n
            \"samples\": [\n              0.9183673469387755, \"n
              0.9453612198959493, \"          0.9723550928531233\n              ], \"semantic_type\": \"\", \"description\": \"\"\n            }, \"column\": \"f1-score\", \"properties\":
            {\n              \"dtype\": \"number\", \"std\":
              0.3926689807140676, \"min\": 0.102272727272728, \"n
              \"max\": 0.985913483827253, \"num_unique_values\": 5, \"n
              \"samples\": [\n                0.102272727272728, \"n
                0.9843932283210849, \"          0.9722622098943156\n                ], \"semantic_type\": \"\", \"description\": \"\"\n              }, \"column\": \"support\", \"properties\":
              {\n                \"dtype\": \"number\", \"std\":
              31154.419955907684, \"min\": 0.9722622098943156, \"n
              \"max\": 56962.0, \"num_unique_values\": 4, \"n
              \"samples\": [\n                98.0, \"          56962.0, \"n
                56864.0\n                ], \"semantic_type\": \"\", \"n
              \"description\": \"\"\n              }\n            }\n          ], \"type\": \"dataframe\"}

```



Logistic Regression AUC-ROC: 0.9720, PR AUC: 0.7622 (Saved internally)

KNN

```
knn_smote_path = "/content/drive/My Drive/Colab  
Notebooks/knn_SMOTE_metrics.json"  
  
if os.path.exists(knn_smote_path):  
    with open(knn_smote_path, 'r') as f:  
        saved_report = json.load(f)  
        print("KNN Classification Report:")  
        display(pd.DataFrame(saved_report).transpose())  
  
    knn_model = KNeighborsClassifier()  
    knn_model.fit(X_train_smote, y_train_smote)  
    y_pred_knn = knn_model.predict(X_test_time)  
    y_proba_knn = knn_model.predict_proba(X_test_time)  
    save_metrics_and_outputs("KNN", y_test, y_pred_knn, y_proba_knn,  
knn_smote_path, internal_auc_scores)  
else:
```

```

print("No saved report found. Training KNN on SMOTE data...")
knn_model = KNeighborsClassifier()
knn_model.fit(X_train_smote, y_train_smote)
y_pred_knn = knn_model.predict(X_test_time)
y_proba_knn = knn_model.predict_proba(X_test_time)
save_metrics_and_outputs("KNN", y_test, y_pred_knn, y_proba_knn,
knn_smote_path, internal_auc_scores)

```

No saved report found. Training KNN on SMOTE data...

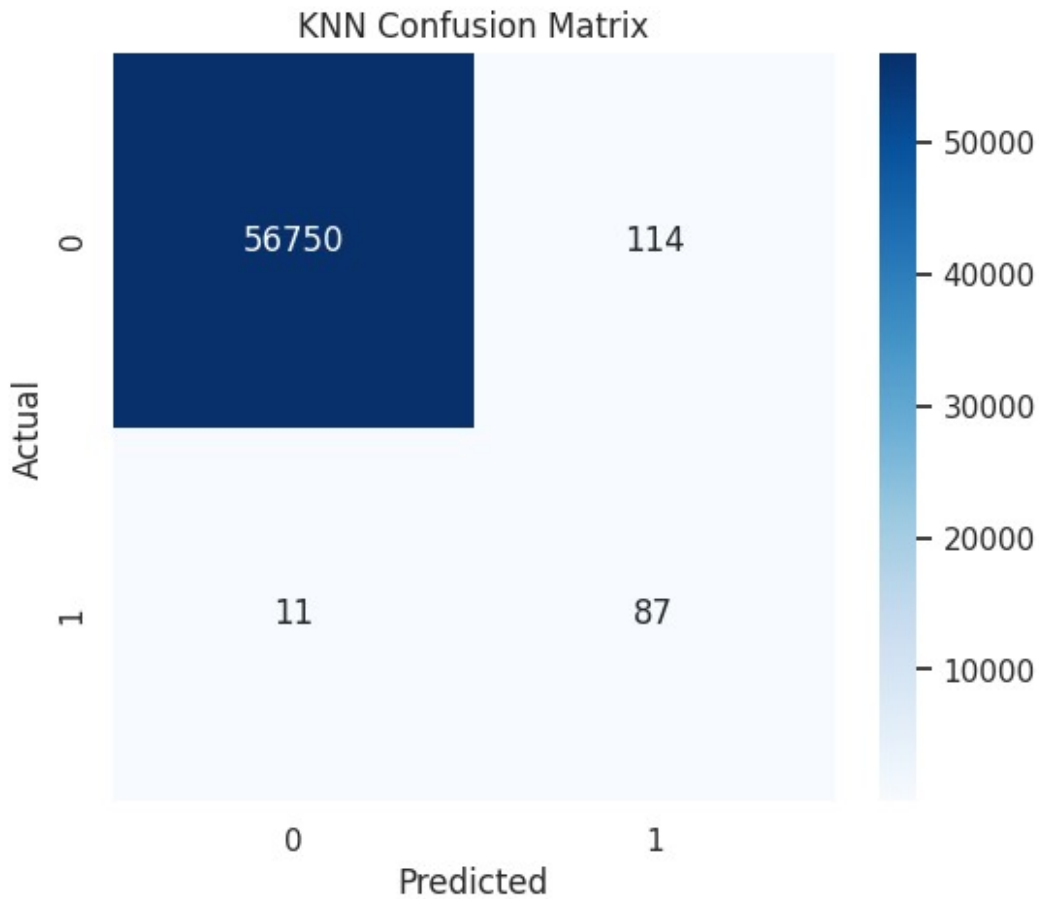
KNN Classification Report:

```

{"summary":{"\n  \"name\": \"      save_metrics_and_outputs(\\\"\\\"KNN\\\"\\\"\",
y_test, y_pred_knn, y_proba_knn, knn_smote_path,
internal_auc_scores)\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"precision\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.25305883813054986, \n        \"min\": 0.43283582089552236, \n        \"max\": 0.9998062049646764, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.43283582089552236, \n          0.9988307634837106, \n          0.9978055545802464\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"recall\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.0492373319231268, \n        \"min\": 0.8877551020408163, \n        \"max\": 0.9979952166572875, \n        \"num_unique_values\": 4, \n        \"samples\": [\n          0.8877551020408163, \n          0.9428751593490519, \n          0.9979952166572875\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"f1-score\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.18616685219859616, \n        \"min\": 0.5819397993311036, \n        \"max\": 0.9988998899889989, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.5819397993311036, \n          0.9981825329986461, \n          0.9978055545802464\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"support\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 31154.412950641894, \n        \"min\": 0.9978055545802464, \n        \"max\": 56962.0, \n        \"num_unique_values\": 4, \n        \"samples\": [\n          98.0, \n          56864.0\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"
    }
  ]\n}, \"type\": \"dataframe\"}

```





KNN AUC-ROC: 0.9484, PR AUC: 0.7570 (Saved internally)

SVM

```
svm_smote_path = "/content/drive/My Drive/Colab  
Notebooks/svm_SMOTE_metrics.json"  
  
if os.path.exists(svm_smote_path):  
    with open(svm_smote_path, 'r') as f:  
        saved_report = json.load(f)  
        print("SVM Classification Report:")  
        display(pd.DataFrame(saved_report).transpose())  
  
    svm_model = SVC(probability=True, random_state=42)  
    svm_model.fit(X_train_smote, y_train_smote)  
    y_pred_svm = svm_model.predict(X_test_time)  
    y_proba_svm = svm_model.predict_proba(X_test_time)  
    save_metrics_and_outputs("SVM", y_test, y_pred_svm, y_proba_svm,  
svm_smote_path, internal_auc_scores)  
else:
```

```

print("No saved report found. Training SVM on SMOTE data...")
svm_model = SVC(probability=True, random_state=42)
svm_model.fit(X_train_smote, y_train_smote)
y_pred_svm = svm_model.predict(X_test_time)
y_proba_svm = svm_model.predict_proba(X_test_time)
save_metrics_and_outputs("SVM", y_test, y_pred_svm, y_proba_svm,
svm_smote_path, internal_auc_scores)

```


No saved report found. Training SVM on SMOTE data...

SVM Classification Report:

```

{"summary":{"\n  \"name\": \"      save_metrics_and_outputs(\\\"\\\"\\\"SVM\\\"\\\"\\\",
y_test, y_pred_svm, y_proba_svm, svm_smote_path,
internal_auc_scores)\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"precision\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.40436014088181743, \n        \"min\": 0.08945686900958466, \n        \"max\": 0.9997501026364172, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.08945686900958466, \n          0.9981839930037598, \n          0.9847442154418735\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"recall\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.057089728278576994, \n        \"min\": 0.8571428571428571, \n        \"max\": 0.9849641249296567, \n        \"num_unique_values\": 4, \n        \"samples\": [\n          0.8571428571428571, \n          0.9210534910362569, \n          0.9849641249296567\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"f1-score\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.3698241902127021, \n        \"min\": 0.16200578592092574, \n        \"max\": 0.9923020365498242, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.16200578592092574, \n          0.990873557343307, \n          0.9847442154418735\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      \"column\": \"support\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 31154.41653271562, \n        \"min\": 0.9847442154418735, \n        \"max\": 56962.0, \n        \"num_unique_values\": 4, \n        \"samples\": [\n          98.0, \n          56962.0, \n          56864.0\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    ]\n  ]\n}, \"type\": \"dataframe\"}

```



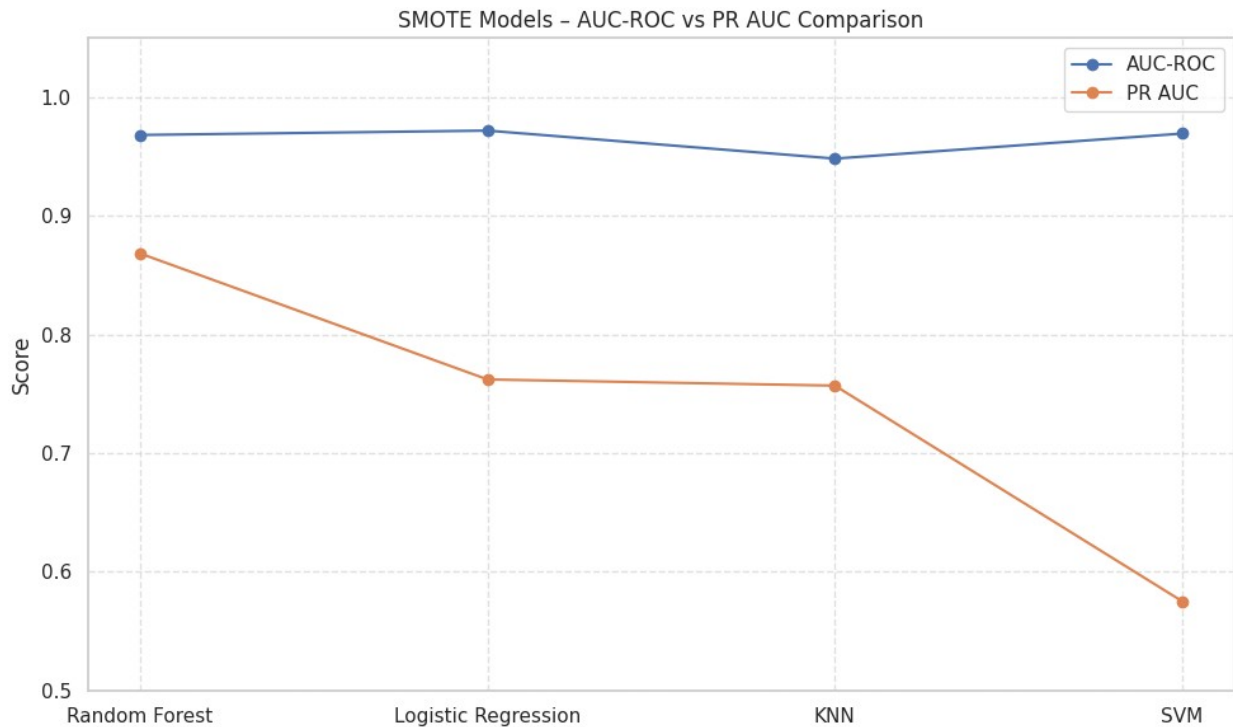
SVM AUC-ROC: 0.9695, PR AUC: 0.5752 (Saved internally)

```
import matplotlib.pyplot as plt

# Model names and manually provided scores
models = ['Random Forest', 'Logistic Regression', 'KNN', 'SVM']
auc_roc_scores = [0.9683, 0.9720, 0.9484, 0.9695]
pr_auc_scores = [0.8684, 0.7622, 0.7570, 0.5752]

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(models, auc_roc_scores, marker='o', label='AUC-ROC')
plt.plot(models, pr_auc_scores, marker='o', label='PR AUC')

# Labels and title
plt.title('SMOTE Models – AUC-ROC vs PR AUC Comparison')
plt.ylabel('Score')
plt.ylim(0.5, 1.05)
plt.grid(True, linestyle='--', alpha=0.6)
plt.legend()
plt.tight_layout()
plt.show()
```



## ADASYN

Random Forest

```
model_name = "Random Forest (ADASYN)"
rf_adasyn_path = "/content/drive/My Drive/Colab
Notebooks/rf_ADASYN_metrics.json"

if os.path.exists(rf_adasyn_path):
    with open(rf_adasyn_path, 'r') as f:
        saved_report = json.load(f)
    print(f"{model_name} Classification Report:")
    display(pd.DataFrame(saved_report).transpose())

    rf_model = RandomForestClassifier(random_state=42)
    rf_model.fit(X_train_adasyn, y_train_adasyn)
    y_pred_rf = rf_model.predict(X_test_time)
    y_proba_rf = rf_model.predict_proba(X_test_time)

    save_metrics_and_outputs(model_name, y_test, y_pred_rf,
y_proba_rf, rf_adasyn_path, internal_auc_scores)
else:
    print(f"No saved report found. Training {model_name} on ADASYN
data...")
    rf_model = RandomForestClassifier(random_state=42)
    rf_model.fit(X_train_adasyn, y_train_adasyn)
```

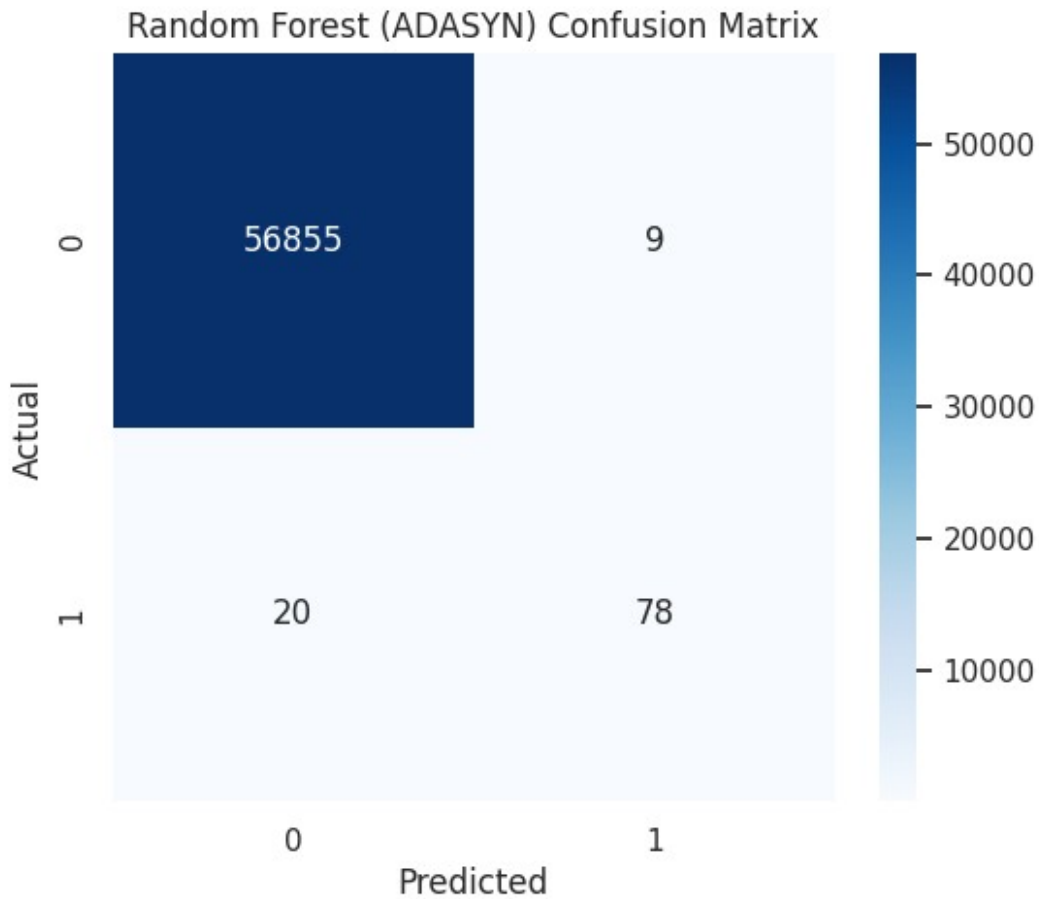
```
y_pred_rf = rf_model.predict(X_test_time)
y_proba_rf = rf_model.predict_proba(X_test_time)
```

```
save_metrics_and_outputs(model_name, y_test, y_pred_rf,
y_proba_rf, rf_adasyn_path, internal_auc_scores)
```

No saved report found. Training Random Forest (ADASYN) on ADASYN data...

#### Random Forest (ADASYN) Classification Report:

```
{"summary":{"\n  \"name\": \"      save_metrics_and_outputs(model_name,
y_test, y_pred_rf, y_proba_rf, rf_adasyn_path,
internal_auc_scores)\" ,\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"precision\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.04605011752147693, \n        \"min\": 0.896551724137931, \n        \"max\": 0.9996483516483516, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.896551724137931, \n          0.9994709795494783, \n          0.9994908886626171\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      }, \n      {\n        \"column\": \"recall\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 0.09107975088840796, \n          \"min\": 0.7959183673469388, \n          \"max\": 0.9998417276308385, \n          \"num_unique_values\": 4, \n          \"samples\": [\n            0.7959183673469388, \n            0.8978800474888886, \n            0.9998417276308385\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n        }, \n        {\n          \"column\": \"f1-score\", \n          \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.06990204303580122, \n            \"min\": 0.8432432432432433, \n            \"max\": 0.9997450302886433, \n            \"num_unique_values\": 5, \n            \"samples\": [\n              0.8432432432432433, \n              0.9994757775389074, \n              0.9994908886626171\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\" \n          }, \n          {\n            \"column\": \"support\", \n            \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 31154.412488438877, \n              \"min\": 0.9994908886626171, \n              \"max\": 56962.0, \n              \"num_unique_values\": 4, \n              \"samples\": [\n                98.0, \n                56962.0, \n                56864.0\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\" \n            } \n          } \n        ] \n      }, \n      \"type\": \"dataframe\"}
```



Random Forest (ADASYN) AUC-ROC: 0.9746, PR AUC: 0.8606 (Saved internally)

Logistic regression

```
model_name = "Logistic Regression (ADASYN)"
lr_adasyn_path = "/content/drive/My Drive/Colab
Notebooks/lr_ADASYN_metrics.json"

if os.path.exists(lr_adasyn_path):
    with open(lr_adasyn_path, 'r') as f:
        saved_report = json.load(f)
    print(f"{model_name} Classification Report:")
    display(pd.DataFrame(saved_report).transpose())

    lr_model = LogisticRegression(max_iter=1000)
    lr_model.fit(X_train_adasyn, y_train_adasyn)
    y_pred_lr = lr_model.predict(X_test_time)
    y_proba_lr = lr_model.predict_proba(X_test_time)

    save_metrics_and_outputs(model_name, y_test, y_pred_lr,
```

```
y_proba_lr, lr_adasyn_path, internal_auc_scores)
```

```
else:
```

```
    print(f"No saved report found. Training {model_name} on ADASYN  
data...")
```

```
    lr_model = LogisticRegression(max_iter=1000)
```

```
    lr_model.fit(X_train_adasyn, y_train_adasyn)
```

```
    y_pred_lr = lr_model.predict(X_test_time)
```

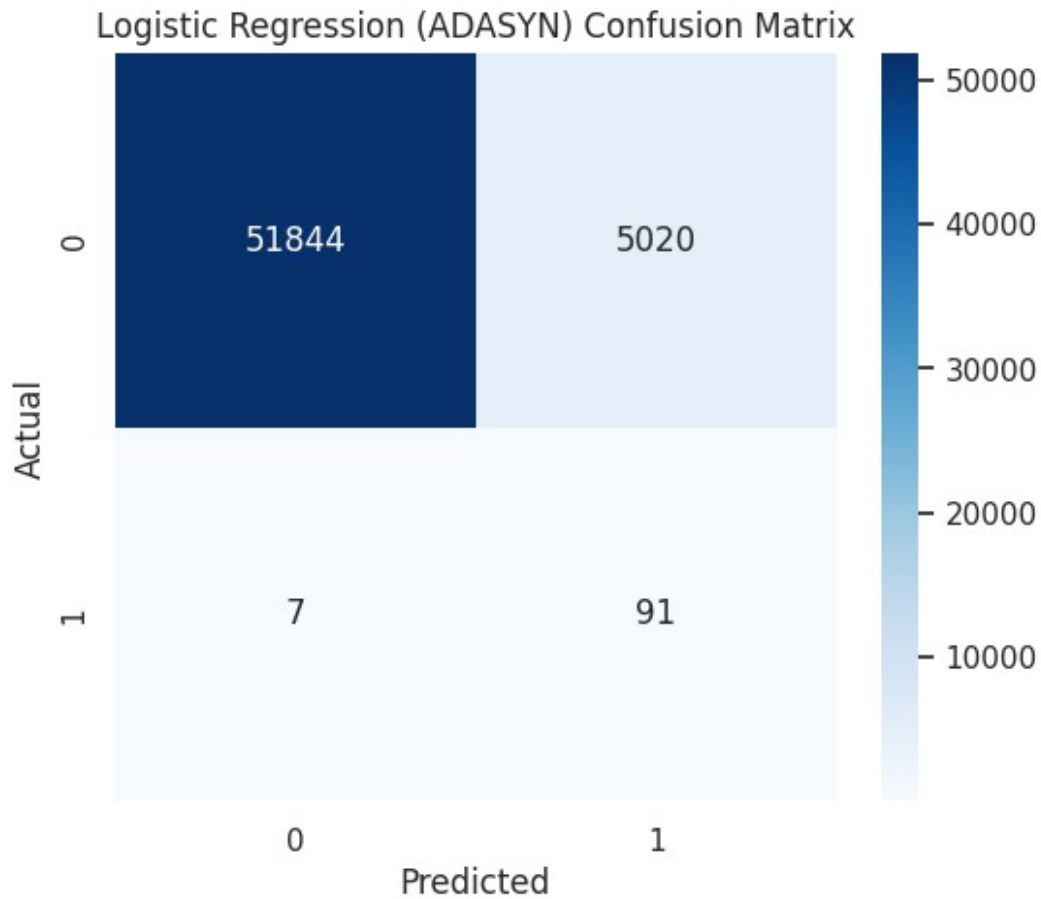
```
    y_proba_lr = lr_model.predict_proba(X_test_time)
```

```
    save_metrics_and_outputs(model_name, y_test, y_pred_lr,  
y_proba_lr, lr_adasyn_path, internal_auc_scores)
```

```
No saved report found. Training Logistic Regression (ADASYN) on ADASYN  
data...
```

```
Logistic Regression (ADASYN) Classification Report:
```

```
{"summary":{"\n  \"name\": \"      save_metrics_and_outputs(model_name,  
y_test, y_pred_lr, y_proba_lr, lr_adasyn_path,  
internal_auc_scores)\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"precision\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.42567332029439975, \n        \"min\": 0.01780473488554099, \n        \"max\": 0.9998649977821065, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.01780473488554099, \n          0.9981754169077716, \n          0.9117481829991925\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\", \n        \"column\": \"precision\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 0.007526836166054648, \n          \"min\": 0.9117191896454699, \n          \"max\": 0.9285714285714286, \n          \"num_unique_values\": 4, \n          \"samples\": [\n            0.9285714285714286, \n            0.9201453091084493, \n            0.9117191896454699\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\", \n          \"column\": \"f1-score\", \n          \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.40396158611943783, \n            \"min\": 0.03493952774044922, \n            \"max\": 0.9537598307501265, \n            \"num_unique_values\": 5, \n            \"samples\": [\n              0.03493952774044922, \n              0.9521790507618019, \n              0.9117481829991925\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\", \n            \"column\": \"support\", \n            \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 31154.436551898303, \n              \"min\": 0.9117481829991925, \n              \"max\": 56962.0, \n              \"num_unique_values\": 4, \n              \"samples\": [\n                98.0, \n                56962.0, \n                56864.0\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\" \n            } \n          } \n        ] \n      }, \n      \"type\": \"dataframe\"}
```



Logistic Regression (ADASYN) AUC-ROC: 0.9739, PR AUC: 0.7670 (Saved internally)

KNN

```
model_name = "KNN (ADASYN)"
knn_adasyn_path = "/content/drive/My Drive/Colab
Notebooks/knn_ADASYN_metrics.json"

if os.path.exists(knn_adasyn_path):
    with open(knn_adasyn_path, 'r') as f:
        saved_report = json.load(f)
    print(f"{model_name} Classification Report:")
    display(pd.DataFrame(saved_report).transpose())

    knn_model = KNeighborsClassifier()
    knn_model.fit(X_train_adasyn, y_train_adasyn)
    y_pred_knn = knn_model.predict(X_test_time)
    y_proba_knn = knn_model.predict_proba(X_test_time)

    save_metrics_and_outputs(model_name, y_test, y_pred_knn,
```



```

y_proba_knn, knn_adasyn_path, internal_auc_scores)
else:
    print(f"No saved report found. Training {model_name} on ADASYN
data...")
    knn_model = KNeighborsClassifier()
    knn_model.fit(X_train_adasyn, y_train_adasyn)
    y_pred_knn = knn_model.predict(X_test_time)
    y_proba_knn = knn_model.predict_proba(X_test_time)

    save_metrics_and_outputs(model_name, y_test, y_pred_knn,
y_proba_knn, knn_adasyn_path, internal_auc_scores)

```

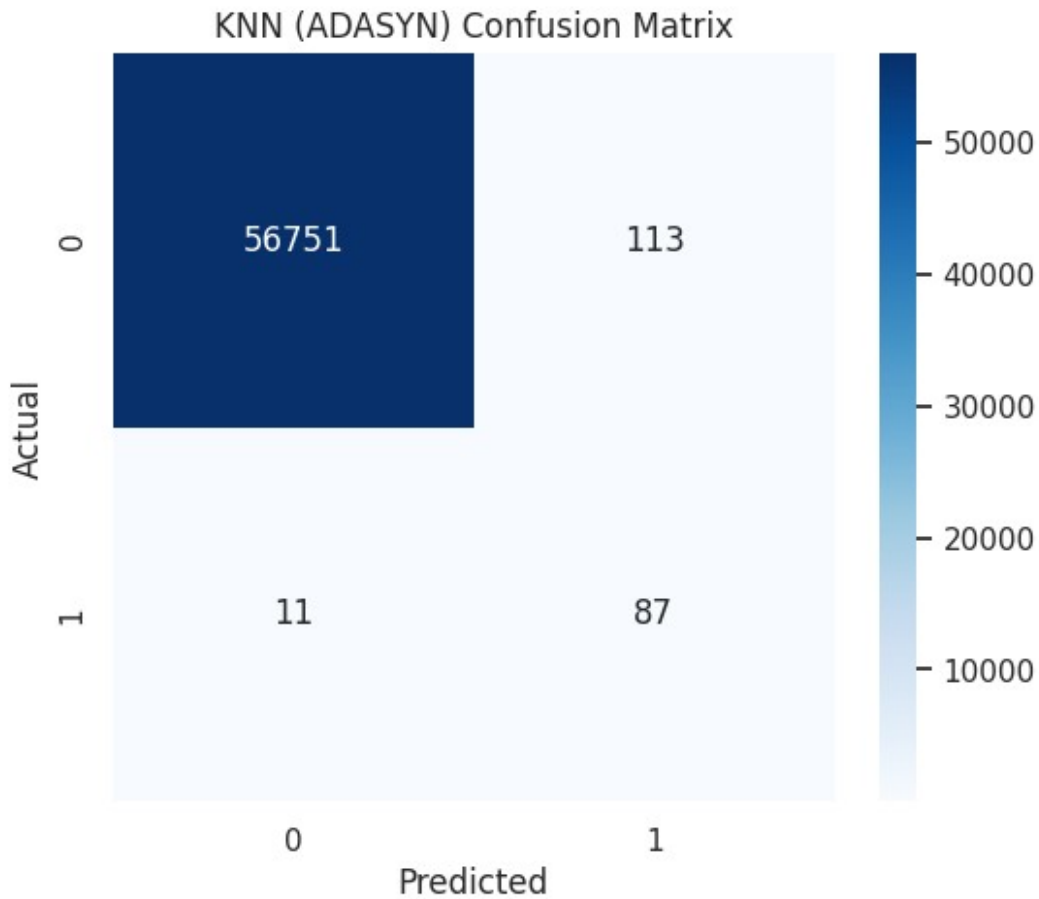
No saved report found. Training KNN (ADASYN) on ADASYN data...

#### KNN (ADASYN) Classification Report:

```

{"summary":{"\n  \"name\": \"      save_metrics_and_outputs(model_name,
y_test, y_pred_knn, y_proba_knn, knn_adasyn_path,
internal_auc_scores)\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"precision\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.2520945427647179, \n        \"min\": 0.435, \n        \"max\": 0.999806208378845, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.435, \n          0.9988344902435772, \n          0.9978231101436045 \n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }, \n      {\n        \"column\": \"recall\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 0.049245186403323926, \n          \"min\": 0.8877551020408163, \n          \"max\": 0.9980128024760833, \n          \"num_unique_values\": 4, \n          \"samples\": [\n            0.8877551020408163, \n            0.9428839522584498, \n            0.9980128024760833 \n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n        }, \n        {\n          \"column\": \"f1-score\", \n          \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.18529948832050916, \n            \"min\": 0.5838926174496645, \n            \"max\": 0.9989087004734832, \n            \"num_unique_values\": 5, \n            \"samples\": [\n              0.5838926174496645, \n              0.9981946880417509, \n              0.9978231101436045 \n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\"\n          }, \n          {\n            \"column\": \"support\", \n            \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 31154.412945827276, \n              \"min\": 0.9978231101436045, \n              \"max\": 56962.0, \n              \"num_unique_values\": 4, \n              \"samples\": [\n                98.0, \n                56962.0, \n                56864.0 \n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\"\n            } \n          ] \n        } \n      ], \n      \"type\": \"dataframe\"}

```



KNN (ADASYN) AUC-ROC: 0.9485, PR AUC: 0.7618 (Saved internally)

SVM

```
model_name = "SVM (ADASYN)"
svm_adasyn_path = "/content/drive/My Drive/Colab
Notebooks/svm_ADASYN_metrics.json"

if os.path.exists(svm_adasyn_path):
    with open(svm_adasyn_path, 'r') as f:
        saved_report = json.load(f)
    print(f"{model_name} Classification Report:")
    display(pd.DataFrame(saved_report).transpose())

    svm_model = SVC(probability=True)
    svm_model.fit(X_train_adasyn, y_train_adasyn)
    y_pred_svm = svm_model.predict(X_test_time)
    y_proba_svm = svm_model.predict_proba(X_test_time)

    save_metrics_and_outputs(model_name, y_test, y_pred_svm,
y_proba_svm, svm_adasyn_path, internal_auc_scores)
```

```

else:
    print(f"No saved report found. Training {model_name} on ADASYN
data...")
    svm_model = SVC(probability=True)
    svm_model.fit(X_train_adasyn, y_train_adasyn)
    y_pred_svm = svm_model.predict(X_test_time)
    y_proba_svm = svm_model.predict_proba(X_test_time)

    save_metrics_and_outputs(model_name, y_test, y_pred_svm,
y_proba_svm, svm_adasyn_path, internal_auc_scores)

```


No saved report found. Training SVM (ADASYN) on ADASYN data...

SVM (ADASYN) Classification Report:

```

{"summary":{"\n  \"name\": \"      save_metrics_and_outputs(model_name,
y_test, y_pred_svm, y_proba_svm, svm_adasyn_path,
internal_auc_scores)\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"precision\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 0.40377009128828484, \n        \"min\": 0.09135802469135802, \n        \"max\": 0.9995725886878473, \n        \"num_unique_values\": 5, \n        \"samples\": [\n          0.09135802469135802, \n          0.9980100552923967, \n          0.9866577718478986\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\", \n        \"column\": \"precision\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 0.10359963202712717, \n          \"min\": 0.7551020408163265, \n          \"max\": 0.9870568373663478, \n          \"num_unique_values\": 4, \n          \"samples\": [\n            0.7551020408163265, \n            0.8710794390913371, \n            0.9870568373663478\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\", \n          \"column\": \"f1-score\", \n          \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.3699716672858355, \n            \"min\": 0.16299559471365638, \n            \"max\": 0.9932752884547321, \n            \"num_unique_values\": 5, \n            \"samples\": [\n              0.16299559471365638, \n              0.9918468377334332, \n              0.9866577718478986\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\", \n            \"column\": \"support\", \n            \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 31154.416007922515, \n              \"min\": 0.9866577718478986, \n              \"max\": 56962.0, \n              \"num_unique_values\": 4, \n              \"samples\": [\n                98.0, \n                56962.0, \n                56864.0\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\" \n            } \n          } \n        ] \n      } \n    } \n  ], \n  \"type\": \"dataframe\"}

```



SVM (ADASYN) AUC-ROC: 0.9254, PR AUC: 0.4090 (Saved internally)

```
import matplotlib.pyplot as plt
```

```
# PR AUC scores
```

```
models = [
```

```
    'Random Forest (SMOTE)', 'Logistic Regression (SMOTE)',
```

```
    'KNN (SMOTE)', 'SVM (SMOTE)',
```

```
    'Random Forest (ADASYN)', 'Logistic Regression (ADASYN)',
```

```
    'KNN (ADASYN)', 'SVM (ADASYN)'
```

```
]
```

```
pr_auc_scores = [0.8684, 0.7622, 0.7570, 0.5752, 0.8606, 0.7670,
0.7618, 0.4090]
```

```
# Plotting
```

```
plt.figure(figsize=(12, 6))
```

```
bars = plt.bar(models, pr_auc_scores, color='skyblue')
```

```
bars[0].set_color('green') # Highlight best performing (Random Forest SMOTE)
```

```
# Annotate scores
```

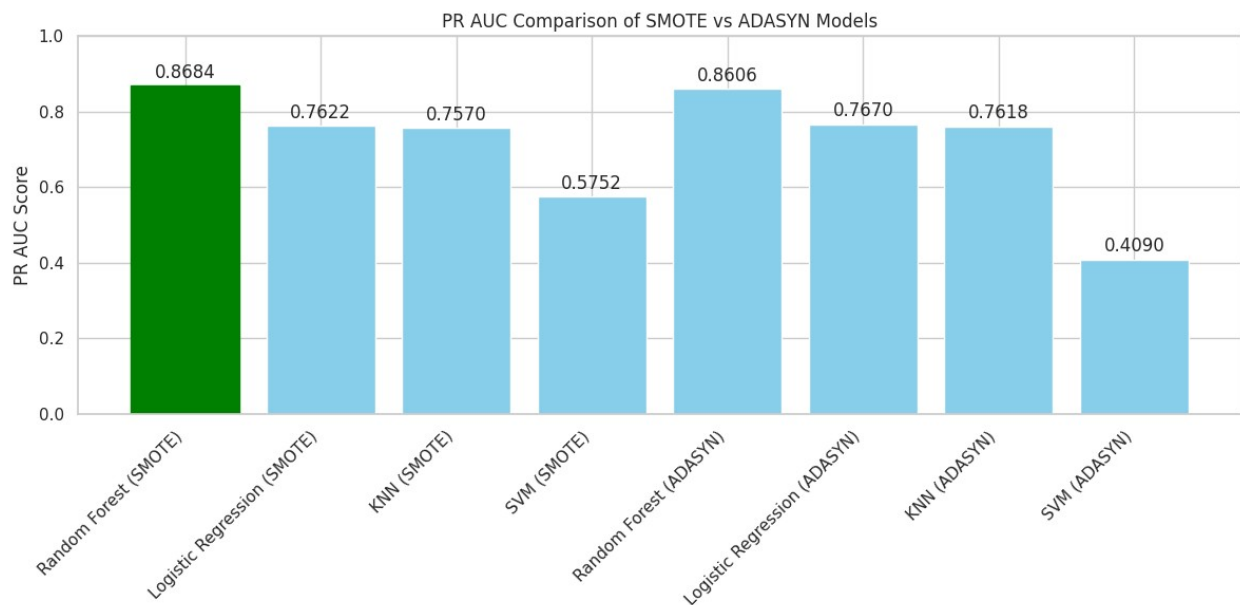
```
for bar in bars:
```

```

        yval = bar.get_height()
        plt.text(bar.get_x() + bar.get_width()/2.0, yval + 0.01,
f'{yval:.4f}', ha='center', va='bottom')

plt.title('PR AUC Comparison of SMOTE vs ADASYN Models')
plt.ylabel('PR AUC Score')
plt.ylim(0, 1)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

```



```

import joblib

# === Finalize and Save the Model ===
final_rf_model = RandomForestClassifier(random_state=42)
final_rf_model.fit(X_train_smote, y_train_smote)

# Save the trained model
joblib.dump(final_rf_model, "final_rf_smote_thresh050_model.pkl")
print(" Final Random Forest model (SMOTE, threshold=0.50) saved
successfully.")

# Optional: Save your threshold value separately (for consistent
inference)
with open("rf_threshold.txt", "w") as f:
    f.write("0.50")

Final Random Forest model (SMOTE, threshold=0.50) saved successfully.

```